

1. The first thought is that I'll be curious how well you're able to scrape news articles - whether you can get sufficient numbers to make the analysis interesting. I'm not sure how many will be behind paywalls, etc. If you get stuck here, please reach out and we can figure out if there are other strategies. I know, for example, that there is a free tier for Diffbot (<https://www.diffbot.com/data/article/>) that may be simpler for gathering news articles.
2. My other thought is that you may be able to do some interesting analysis using LLMs. It is tricky to get calibrated output from LLMs, but I'm happy to discuss how you might structure this. I'd have to look at whether any providers currently offer free (or free for students, or just really cheap) API usage. An alternative may be using a smaller language model locally, which can be feasible, as well.
3. The biggest challenges would probably be the availability of data: some social media sites and news organizations are behind paywalls or logins. Another challenge may be how to handle multiple languages. I know there is probably a lot of international news that could be in any language or and local news would presumably be in Nepali. Finding information about how locals feel might differ than for how the news was covered internationally.
4. The topic is an interesting one. I think the question of how people reacted to a Discord election is also fascinating. It is understandable that many news articles may be behind paywalls and the approach of combining links and news stories is a way of handling the relative lack of English language data.
5. I think it would have been useful to show some specific examples of either strong positive or strong negative sentiment in the articles you scraped. Many people may not be familiar with the election, so this would give some sense of the range of reactions.
6. In terms of methodology, you provided brief explanations of how each method works, but I think that the fact that they gave "opposite" results should've encouraged a bit more digging into the differences - or hand-labeling of a few examples to have more concrete information on which was performing better. That type of approach might also have led you towards additional steps in the analysis, such as splitting articles into chunks (sentences/paragraphs) and scoring them individually or, at minimum, given more quantitative/qualitative explanation for which is working better. With only a few dozen articles after joining data sets, it's conceivable to have labeled many/most/all of them manually to have more quantitative results.
7. For your final paper, I'd suggest going into more detail on how the 2 different methods work, finding a more concrete way to recommend which works better (even showing examples of articles where they disagree might be insightful) as well as showing a bit more information about the data itself - some examples, some EDA, etc. Another thing that might be useful is taking the word clouds and making them relative, such as the ratio of positive to negative use of a particular word.
8. Since the presentation, you did add some more details with example records and records where VADER and TextBlob disagreed, as well as comparing how things differed on different data pulls. Scraping data can also be more challenging than downloading data, such as you found with the
9. Some of the areas I mentioned before I think are still true:

10. You could've shown more exploratory data analysis. For example, by finding specific statements expressing very positive or negative sentiment (perhaps from manual review - there are only 23 articles) This is also where splitting into chunks, like I mentioned, could let you do more analysis. In a news article, you may have examples of people expressing both positive and negative sentiment, so perhaps the methods would be more consistent with each other if you had first broken into paragraphs or smaller chunks of text.
11. The methods have more detail now on how they work, but it is still pretty "high level".
12. I see you experimented with a BERT model but did not report the results, other than a brief mention.
13. The code is straightforward
14. The paper itself is reasonably clear, but I think you could've tried to focus more on answering the original business question about the sentiment around the election vs the comparison of the two libraries.
15. Create a transformer from scratch to do better analysis.